

[hw2]

September 2, 2024

Download the dataset from: <https://github.com/bellawillrise/Introduction-to-Numerical-Computing-in-Python/>

Submit a pdf file, which is a rendered saved version of the jupyter notebook. Make sure to execute all the codes so the output can be viewed in the pdf.

Also include the link to the public github repository where the jupyter notebook for the assignment is uploaded.

Link to the github repository: <https://github.com/dreeew05/CMSC-197/tree/main/Assignment%201>

```
[49]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[50]: # %matplotlib inline
```

```
[51]: data = pd.read_csv("data/movie_metadata_cleaned.csv")
```

```
[52]: data.head(2)
```

```
[52]: Unnamed: 0      movie_title  color \
0      0      b'Avatar'  Color
1      1  b"Pirates of the Caribbean: At World's End"  Color

      director_name  num_critic_for_reviews  duration  director_facebook_likes \
0  James Cameron      723.0      178.0      0.0
1  Gore Verbinski      302.0      169.0      563.0

      actor_3_facebook_likes  actor_2_name  actor_1_facebook_likes  ... \
0      855.0  Joel David Moore      1000.0  ...
1      1000.0  Orlando Bloom      40000.0  ...

      num_user_for_reviews  language  country  content_rating  budget \
0      3054.0  English  USA  PG-13  237000000.0
1      1238.0  English  USA  PG-13  300000000.0
```

```

title_year  actor_2_facebook_likes  imdb_score  aspect_ratio  \
0      2009.0                936.0         7.9         1.78
1      2007.0                5000.0        7.1         2.35

movie_facebook_likes
0          33000.0
1           0.0

[2 rows x 29 columns]

```

0.1 Get the top 10 directors with most movies directed and use a boxplot for their gross earnings

```

[53]: # Filter data [Remove directors that are named '0']
filtered_directors = data[data['director_name'] != '0']

# Group by director and get the top 10 directors with most movies directed
top_directors = filtered_directors.groupby('director_name').size().
    ↪sort_values(ascending=False).head(10)

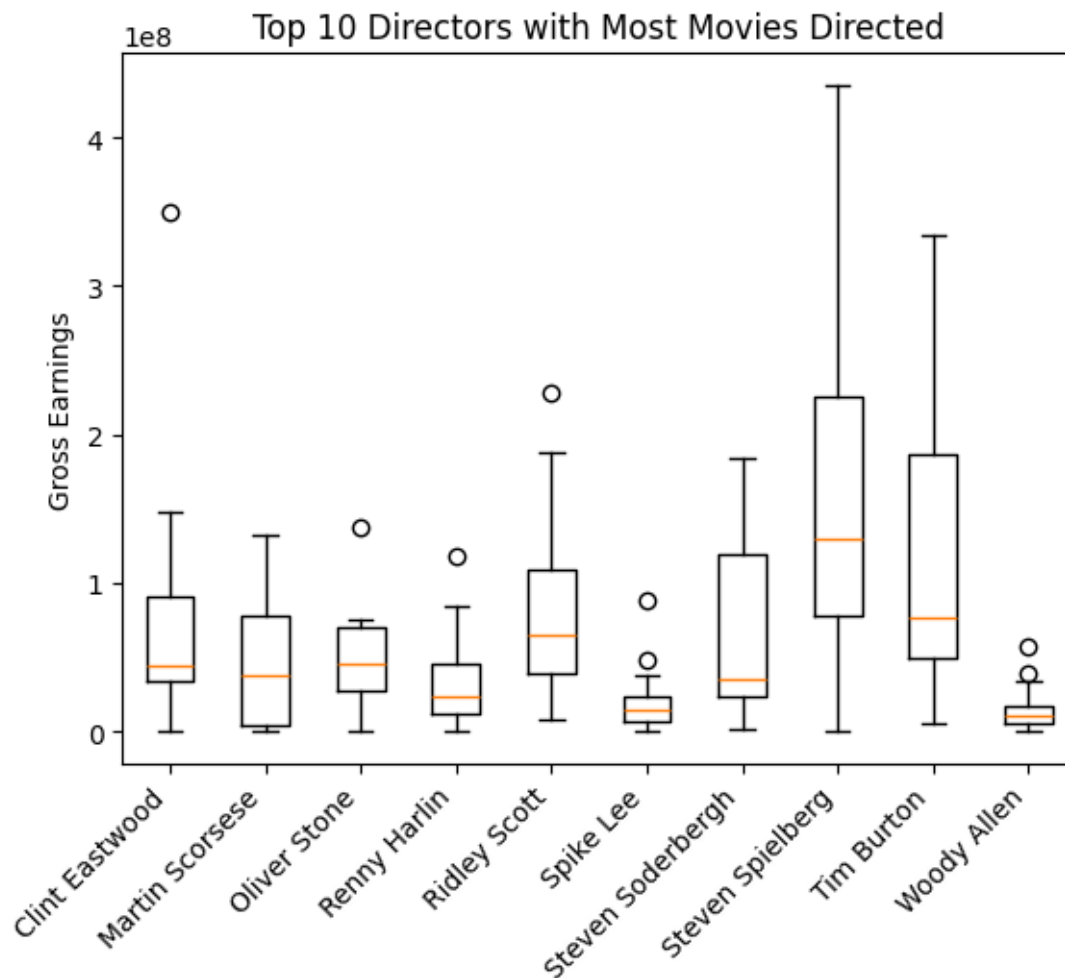
# Filter the original dataframe to include only these top directors
top_directors_data = data[data['director_name'].isin(top_directors.index)]

# Group the data by director and get the gross earnings for each of their movie
gross_earnings_by_director = top_directors_data.
    ↪groupby('director_name')['gross'].apply(list)

# Create a boxplot for the top directors' gross earnings
plt.boxplot(gross_earnings_by_director, tick_labels=gross_earnings_by_director.
    ↪index)

plt.title('Top 10 Directors with Most Movies Directed')
plt.ylabel('Gross Earnings')
plt.xticks(rotation=45, ha="right")
plt.show()

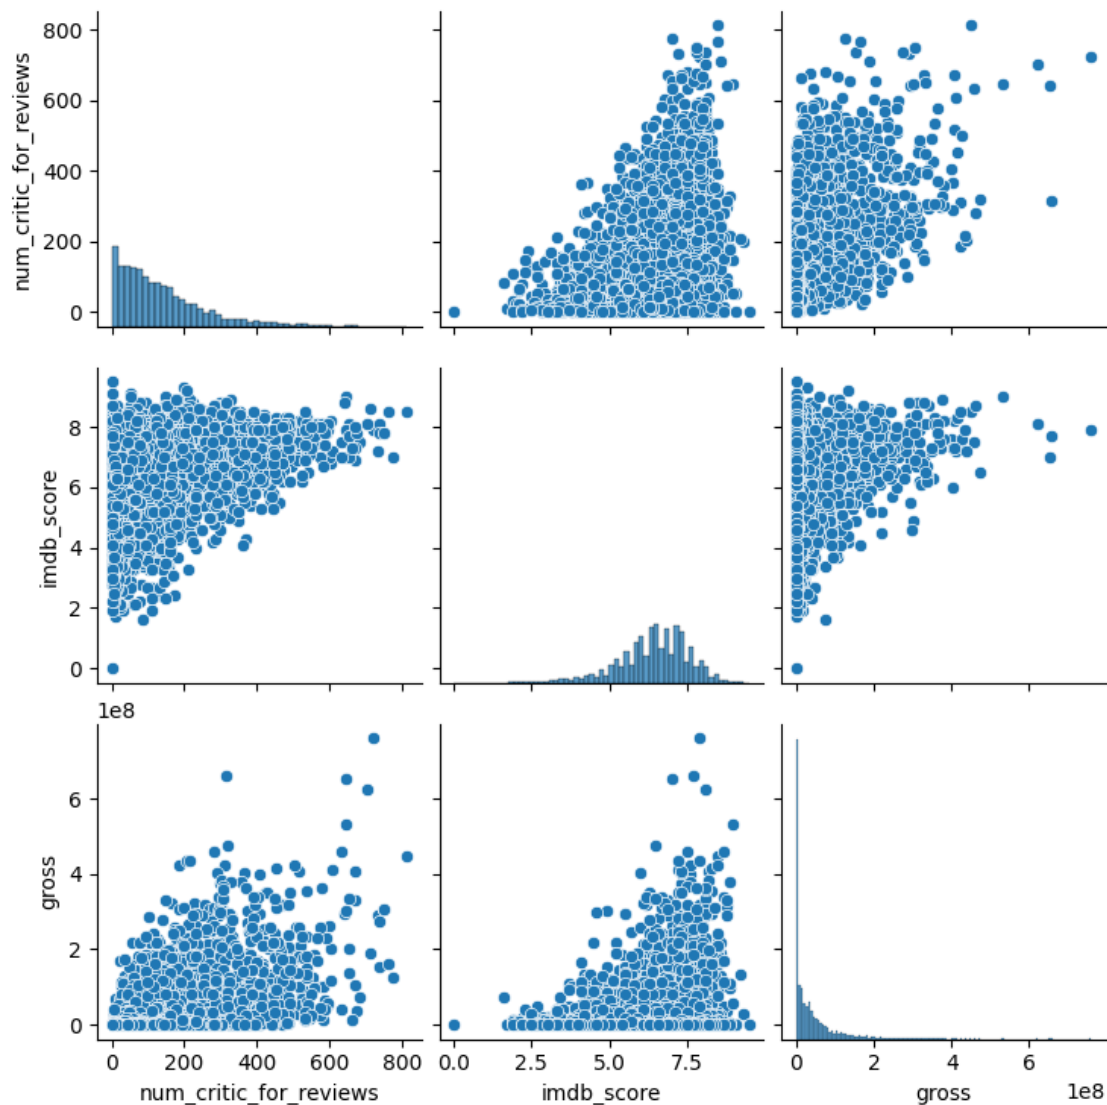
```



0.2 Plot the following variables in one graph:

- num_critic_for_reviews
- IMDB score
- gross

```
[54]: sns.pairplot(data[['num_critic_for_reviews', 'imdb_score', 'gross']])
plt.show()
```



0.3 Compute Sales (Gross - Budget), add it as another column

```
[55]: data['computed_sales'] = data['gross'] - data['budget']
      data.head()
```

```
[55]: Unnamed: 0      movie_title  color \
0      0      b'Avatar'    Color
1      1  b"Pirates of the Caribbean: At World's End"  Color
2      2      b'Spectre'    Color
3      3      b'The Dark Knight Rises'    Color
4      4  b'Star Wars: Episode VII - The Force Awakens ...    0
```

	director_name	num_critic_for_reviews	duration	\
0	James Cameron	723.0	178.0	
1	Gore Verbinski	302.0	169.0	
2	Sam Mendes	602.0	148.0	
3	Christopher Nolan	813.0	164.0	
4	Doug Walker	0.0	0.0	

	director_facebook_likes	actor_3_facebook_likes	actor_2_name	\
0	0.0	855.0	Joel David Moore	
1	563.0	1000.0	Orlando Bloom	
2	0.0	161.0	Rory Kinnear	
3	22000.0	23000.0	Christian Bale	
4	131.0	0.0	Rob Walker	

	actor_1_facebook_likes	...	language	country	content_rating	budget	\
0	1000.0	...	English	USA	PG-13	237000000.0	
1	40000.0	...	English	USA	PG-13	300000000.0	
2	11000.0	...	English	UK	PG-13	245000000.0	
3	27000.0	...	English	USA	PG-13	250000000.0	
4	131.0	...	0	0	0	0.0	

	title_year	actor_2_facebook_likes	imdb_score	aspect_ratio	\
0	2009.0	936.0	7.9	1.78	
1	2007.0	5000.0	7.1	2.35	
2	2015.0	393.0	6.8	2.35	
3	2012.0	23000.0	8.5	2.35	
4	0.0	12.0	7.1	0.00	

	movie_facebook_likes	computed_sales
0	33000.0	523505847.0
1	0.0	9404152.0
2	85000.0	-44925825.0
3	164000.0	198130642.0
4	0.0	0.0

[5 rows x 30 columns]

0.4 Which directors garnered the most total sales?

```
[56]: # filtered_directors is a data frame used above
# Get the total sales using aggregation
director_agg = filtered_directors.groupby('director_name').agg(
    total_sales = ('gross', 'sum')
)
# Sort and print the directors with most total sales
director_agg.sort_values(by='total_sales', ascending=False).head()
```

```
[56]:
```

	total_sales
director_name	
Steven Spielberg	4.114233e+09
Peter Jackson	2.592969e+09
Michael Bay	2.231243e+09
Tim Burton	2.071275e+09
Sam Raimi	2.049549e+09

0.5 Plot sales and average likes as a scatterplot. Fit it with a line.


```
[57]: # Add average_likes as a new column
total_likes = [
    'movie_facebook_likes',
    'actor_1_facebook_likes',
    'actor_2_facebook_likes',
    'actor_3_facebook_likes',
    'director_facebook_likes',
    'cast_total_facebook_likes'
]
data['average_likes'] = data[total_likes].mean(axis=1)

plt.scatter(data['computed_sales'], data['average_likes'], color='blue',
            alpha=0.5)

#Linear Fit
slope, intercept = np.polyfit(data['computed_sales'], data['average_likes'], 1)
fit_line = slope * data['computed_sales'] + intercept

# Plot the fit line
plt.plot(data['computed_sales'], fit_line, color='red', linewidth=2,
        label=f'Fit Line: y = {slope:.2f}x + {intercept:.2f}')

plt.title('Scatter Plot of Sales vs. Average Likes with Fit Line')
plt.xlabel('Sales')
plt.ylabel('Average Likes')
plt.legend()
plt.show()
```



0.6 Which of these genres are the most profitable? Plot their sales using different histograms, superimposed in the same axis.

- Romance
- Comedy
- Action
- Fantasy

```
[58]: # Filter data by the specific genres
romance = data[data['genres'].str.contains('Romance', na=False)]
comedy = data[data['genres'].str.contains('Comedy', na=False)]
action = data[data['genres'].str.contains('Action', na=False)]
fantasy = data[data['genres'].str.contains('Fantasy', na=False)]

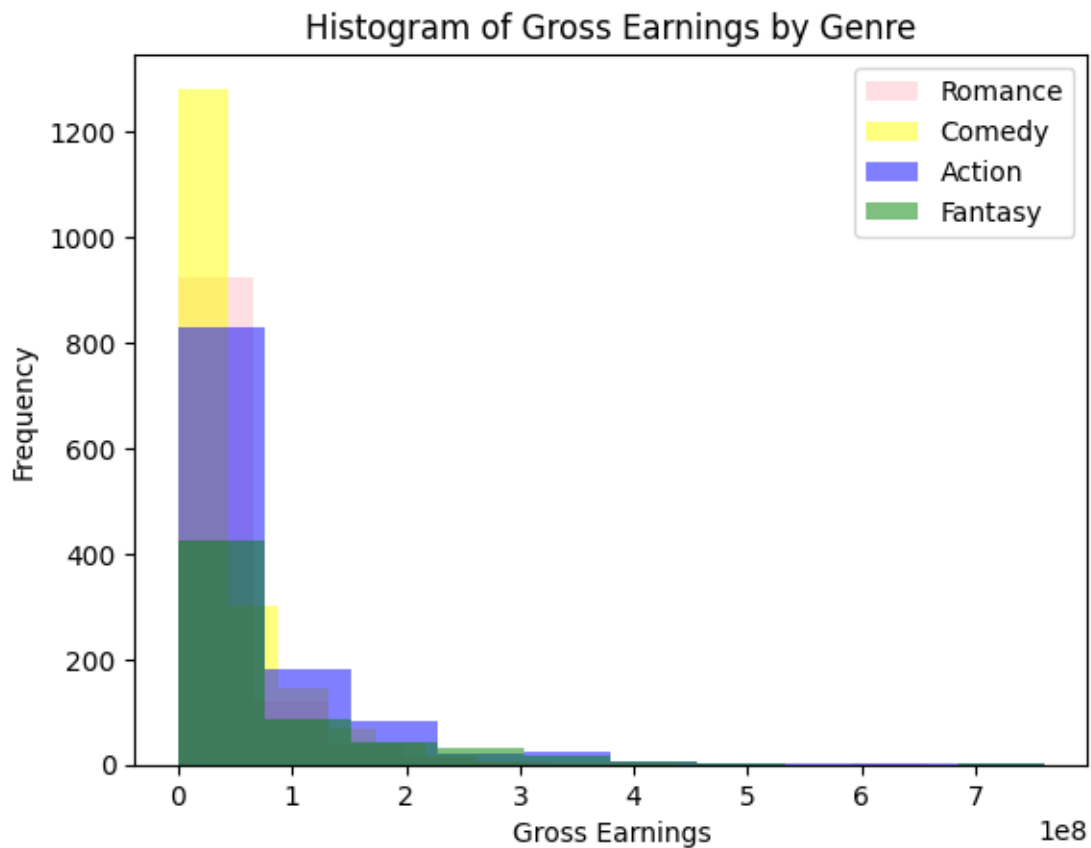
# Romance
plt.hist(romance['gross'].dropna(), alpha=0.5, label='Romance', color='pink')

# Comedy
plt.hist(comedy['gross'].dropna(), alpha=0.5, label='Comedy', color='yellow')
```

```
# Action
plt.hist(action['gross'].dropna(), alpha=0.5, label='Action', color='blue')

# Fantasy
plt.hist(fantasy['gross'].dropna(), alpha=0.5, label='Fantasy', color='green')

plt.xlabel('Gross Earnings')
plt.ylabel('Frequency')
plt.title('Histogram of Gross Earnings by Genre')
plt.legend()
plt.show()
```



0.7 For each of movie, compute average likes of the three actors and store it as a new variable

Read up on the mean function.

Store it as a new column, `average_actor_likes`.

0.8 Copying the whole dataframe

```
[59]: df = data.copy()
df.head()
```

```
[59]: Unnamed: 0      movie_title  color \
0      0      b'Avatar'  Color
1      1  b"Pirates of the Caribbean: At World's End"  Color
2      2      b'Spectre'  Color
3      3  b'The Dark Knight Rises'  Color
4      4  b'Star Wars: Episode VII - The Force Awakens ...    0

      director_name  num_critic_for_reviews  duration \
0      James Cameron      723.0      178.0
1      Gore Verbinski      302.0      169.0
2      Sam Mendes      602.0      148.0
3  Christopher Nolan      813.0      164.0
4      Doug Walker      0.0      0.0

      director_facebook_likes  actor_3_facebook_likes  actor_2_name \
0      0.0      855.0  Joel David Moore
1      563.0      1000.0  Orlando Bloom
2      0.0      161.0  Rory Kinnear
3      22000.0      23000.0  Christian Bale
4      131.0      0.0  Rob Walker

      actor_1_facebook_likes  ...  country  content_rating  budget \
0      1000.0  ...  USA  PG-13  237000000.0
1      40000.0  ...  USA  PG-13  300000000.0
2      11000.0  ...  UK  PG-13  245000000.0
3      27000.0  ...  USA  PG-13  250000000.0
4      131.0  ...  0  0  0.0

      title_year  actor_2_facebook_likes  imdb_score  aspect_ratio \
0      2009.0      936.0      7.9      1.78
1      2007.0      5000.0      7.1      2.35
2      2015.0      393.0      6.8      2.35
3      2012.0      23000.0      8.5      2.35
4      0.0      12.0      7.1      0.00

      movie_facebook_likes  computed_sales  average_likes
0      33000.0      523505847.0      6770.833333
1      0.0      9404152.0      15818.833333
2      85000.0      -44925825.0      18042.333333
3      164000.0      198130642.0      60959.833333
4      0.0      0.0      69.500000
```

[5 rows x 31 columns]

0.9 Min-Max Normalization

Normalization is a technique often applied as part of data preparation for machine learning. The goal of normalization is to change the values of numeric columns in the dataset to a common scale, without distorting differences in the ranges of values. For machine learning, every dataset does not require normalization. It is required only when features have different ranges.

The min-max approach (often called normalization) rescales the feature to a hard and fast range of [0,1] by subtracting the minimum value of the feature then dividing by the range. We can apply the min-max scaling in Pandas using the `.min()` and `.max()` methods.

$$x_{scaled} = \frac{x - x_{min}}{x_{max} - x_{min}}$$

0.9.1 Normalize each numeric column (those that have types integer or float) of the copied dataframe (df)

```
[60]: numeric_columns = df.select_dtypes(include=['float64', 'int64']).columns
df[numeric_columns] = df[numeric_columns].apply(lambda x: (x - x.min()) / (x.
↳max() - x.min()))
df.head()
```

```
[60]: Unnamed: 0      movie_title  color \
0      0.000000      b'Avatar'  Color
1      0.000198  b"Pirates of the Caribbean: At World's End"  Color
2      0.000397      b'Spectre'  Color
3      0.000595  b'The Dark Knight Rises'  Color
4      0.000793  b'Star Wars: Episode VII - The Force Awakens ...    0

      director_name  num_critic_for_reviews  duration \
0      James Cameron      0.889299  0.941799
1      Gore Verbinski      0.371464  0.894180
2      Sam Mendes      0.740467  0.783069
3  Christopher Nolan      1.000000  0.867725
4      Doug Walker      0.000000  0.000000

      director_facebook_likes  actor_3_facebook_likes  actor_2_name \
0      0.000000      0.037174  Joel David Moore
1      0.024478      0.043478  Orlando Bloom
2      0.000000      0.007000  Rory Kinnear
3      0.956522      1.000000  Christian Bale
4      0.005696      0.000000  Rob Walker

      actor_1_facebook_likes  ...  country  content_rating  budget  title_year \
0      0.001563  ...      USA      PG-13  0.019402  0.996528
1      0.062500  ...      USA      PG-13  0.024559  0.995536
```

2	0.017188	...	UK	PG-13	0.020056	0.999504
3	0.042188	...	USA	PG-13	0.020466	0.998016
4	0.000205	...	0	0	0.000000	0.000000

	actor_2_facebook_likes	imdb_score	aspect_ratio	movie_facebook_likes	\
0	0.006832	0.831579	0.111250		0.094556
1	0.036496	0.747368	0.146875		0.000000
2	0.002869	0.715789	0.146875		0.243553
3	0.167883	0.894737	0.146875		0.469914
4	0.000088	0.747368	0.000000		0.000000

	computed_sales	average_likes
0	1.000000	0.030964
1	0.959637	0.072341
2	0.955371	0.082510
3	0.974454	0.278777
4	0.958898	0.000318

[5 rows x 31 columns]